

Infrastructure as Code Support in Graal CI

JSON-Driven IaC Framework with Pulumi & GitLab CI/CD

Yassine DBAICHI

École Hassania des Travaux Publics

Oracle Labs - GraalVM RISQ Team

Academic Year 2024-2025

Agenda

1. Company & Team Context

Oracle, Oracle Labs, GraalVM
RISQ Team mission and focus

2. Problem & Solution Overview

Current challenges, requirements,
and proposed JSON-driven IaC
solution

3. Project Management

Agile methodology,
communication tools, and
development workflow

4. Technical Implementation

Architecture, design patterns,
CI/CD integration, and deployment

5. Results & Validations

Implementation results, testing
outcomes, and performance
metrics

6. Future Enhancements

Current limitations, future work,
and conclusion

Company & Team Context

Oracle Corporation

Global Technology Leader: Multinational computer technology corporation specializing in database software and cloud solutions

Founded: 1977 by Larry Ellison, Bob Miner, and Ed Oates

Headquarters: Austin, Texas, USA (Oracle Headquarters)

Market Position: One of the largest enterprise software companies worldwide

Core Products: Oracle Database, Oracle Cloud Infrastructure (OCI), Enterprise Applications

Revenue: \$50+ billion annually with 430,000+ employees globally

Oracle Labs - Research & Innovation

Mission: Conduct research to advance the state of the art in computer science and systems

Focus Areas:

Programming languages & compilers

Virtual machines & runtime systems

Database systems & optimization

Cloud infrastructure & security

Research Excellence

Oracle Labs bridges cutting-edge academic research with real-world enterprise applications

Output: Open-source projects, academic publications, and production-ready technologies

Collaboration

Strong partnerships with universities and research institutions worldwide

Key Projects: GraalVM, Truffle Framework, SubstrateVM

GraalVM - High-Performance Polyglot Runtime

What is GraalVM? A universal virtual machine for running applications written in JavaScript, Python, Ruby, R, JVM languages, and more

Key Features:

Polyglot capabilities
(multiple languages)

Native Image compilation

Superior performance optimization

Reduced memory footprint



GraalVM RISQ Team

Team Name: RISQ (Release Infrastructure, Systems & Quality)

Mission: Ensure reliability, scalability, and quality of GraalVM infrastructure and development workflows

Responsibilities:

CI/CD pipeline management and optimization

Infrastructure automation and provisioning

Testing infrastructure maintenance

Developer tooling and experience improvements

Tech Stack: GitLab CI/CD, Oracle Cloud Infrastructure, Python, Bash scripting

The RISQ team acts as the backbone supporting GraalVM development, ensuring smooth operations for hundreds of

Project Context & Internship Focus

Team Challenge

GraalVM requires frequent infrastructure provisioning for testing, CI/CD, and development environments across multiple cloud regions

Manual Bottleneck

Infrastructure management was manual, error-prone, and time-consuming - limiting team agility

Project Goal

Design and implement an automated Infrastructure as Code (IaC) solution integrated with Graal CI pipeline

My Role

Lead developer responsible for architecture design, implementation, and GitLab CI/CD integration

Problem & Solution Overview

Current Infrastructure Management Challenges

Manual Console Operations: Team members had to manually create resources through OCI web console - time-consuming and repetitive

Lack of Version Control: No tracking of infrastructure changes, making rollbacks and audits difficult

Configuration Drift: Resources created manually often diverged from documented standards

Scalability Issues: Provisioning multiple similar environments was labor-intensive

Knowledge Silos: Infrastructure knowledge concentrated in few team members

No CI/CD Integration: Infrastructure provisioning was separate from development workflow

Result: Slow deployment cycles, increased error rates, and reduced

Functional Requirements

FR1 - Declarative Configuration: Users define infrastructure in human-readable format (JSON) without writing Python/Terraform code

FR2 - Resource Support: Support for essential OCI resources (VCN, Subnets, Compute Instances, Storage Buckets, Security Lists, Routing)

FR3 - Dependency Management: Automatic resolution of resource dependencies (e.g., subnet requires VCN, instance requires subnet)

FR4 - CI/CD Integration: Seamless integration with GitLab CI/CD pipeline (plan, preview, apply stages)

FR5 - State Management: Track infrastructure state to enable updates and deletions

FR6 - Multi-Environment: Support for development, staging, and production configurations

Non-Functional Requirements

Performance

- Provision infrastructure in <5 minutes
- Support concurrent deployments
- Minimal API calls to OCI

Security

- Credentials stored securely (GitLab secrets)
- No hardcoded sensitive data
- Audit trail via Git commits

Extensibility

- Plugin architecture for new resources
- Support for future cloud providers
- Custom validation rules

IaC Technology Selection: Terraform vs Pulumi

Terraform

Pros:

- Industry standard with mature ecosystem
- HCL declarative language
- Large provider ecosystem
- Strong community support

Cons:

- Custom DSL (HCL) requires learning
- Limited programming logic & abstraction
- State management complexity
- Hard to build dynamic configurations

Pulumi ✓ (Selected)

Pros:

- Use real programming languages (Python)
- Full programmatic control & abstraction
- Built-in state management
- Excellent OCI support
- Enables framework development on top
- Team Python expertise leverage

Why chosen: Perfect fit for building JSON-driven abstraction layer. Pulumi's programmatic approach allows us to create a framework that hides IaC complexity from users.

How does Pulumi work?

1. Write Code

Define infrastructure using Python
- loops, functions, conditionals available

2. Preview

Run 'pulumi preview' to see what will change before applying

3. Deploy

Pulumi SDK translates Python to OCI API calls and provisions resources

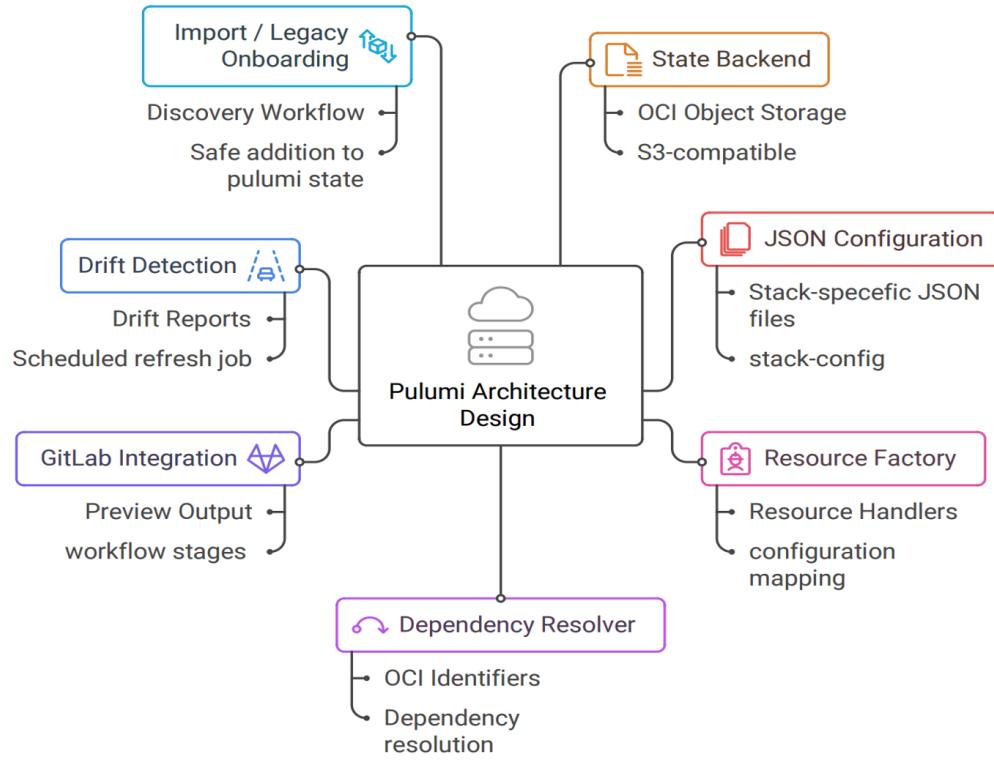
4. State Management

Pulumi tracks all resources in state file (OCI Object Storage)

Python-First: Leverage existing team Python expertise

OCI Provider: First-class Oracle Cloud Infrastructure support

Framework Foundation: Enables building JSON-driven abstraction layer on top



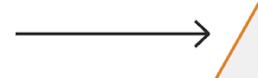
Integration

GitLab CI/CD automates workflow management



Provisioning

Pulumi SDK manages infrastructure resources



Orchestration

Python engine coordinates configuration and execution



Configuration

JSON files define infrastructure requirements



Project Management

Project Management Approach

Methodology: Kanban Agile - visual workflow management with continuous delivery

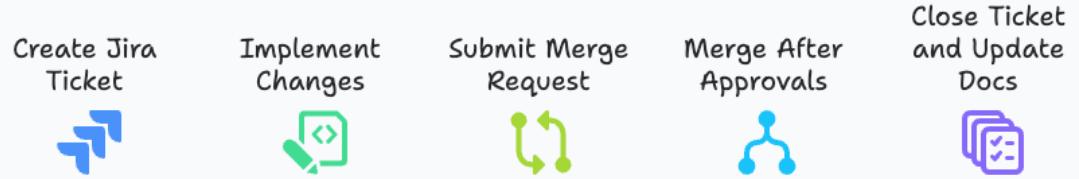
Project Tracker: Task board for tracking work items (To Do, In Progress, Done)

Regular Meetings: Weekly sync meetings with mentor to review progress and blockers

Continuous Flow: Work items pulled as capacity allows, no fixed sprint boundaries

Focus on Delivery: Prioritize completing tasks over starting new ones

Development Task Workflow



Made with  Napkin

Communication & Collaboration Tools



Slack

Daily communication, quick questions, team updates



Zoom

Weekly sync meetings, architecture discussions



Jira

Task tracking, sprint planning, progress monitoring



Outlook

Meeting scheduling, formal communications



Confluence

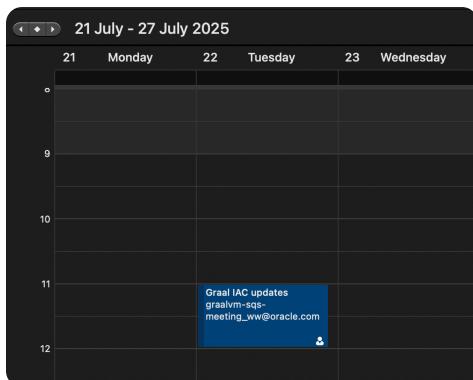
Documentation, design specs, knowledge base



GitLab

Code repository, CI/CD, merge requests

Mentor Check-ins & Team Meetings



Weekly Sync Meetings

Frequency: Every Tuesday

11:00 AM

Duration: 45-60 minutes

Agenda: Progress review, blockers discussion, next steps planning

Participants: Intern + mentor (Hamza Ghaissi)

Ad-hoc Sessions

Technical architecture discussions

Code reviews and pair programming

Debugging complex issues

Design decision consultations

Communication Style:

Open-door policy via Slack - quick responses and collaborative problem-solving

Project Timeline & Key Milestones

Phase 1: Research (2 weeks)

OCI API exploration
Pulumi SDK study
Requirements gathering
Architecture design

Phase 2: Prototype (2 weeks)

Core framework skeleton
VCN + Instance resources
Factory pattern POC
Basic dependency resolver

Phase 3: Development (4 weeks)

All 7 resource types
JSON configuration loader
Dependency resolution logic
Error handling

Phase 4: CI/CD Integration (2 weeks)

GitLab pipeline config
Plan & deploy stages
Manual

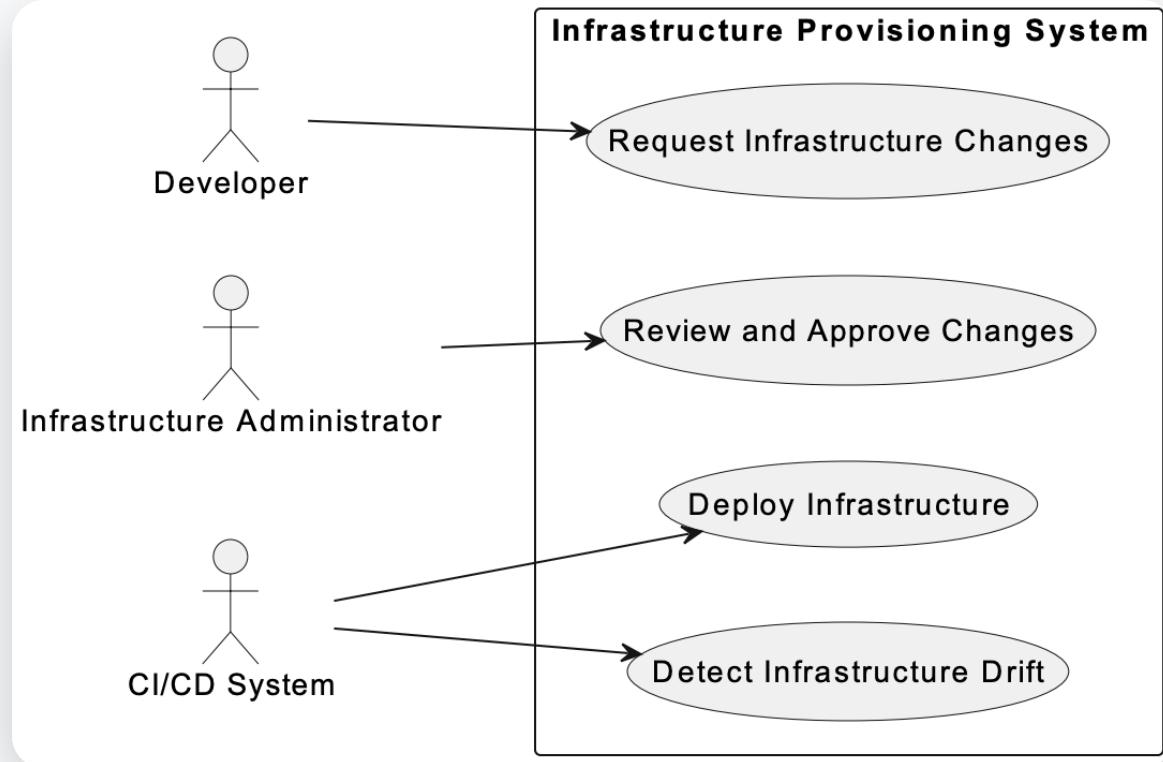
Phase 5: Testing (2 weeks)

Integration testing
Edge case validation
Performance testing

Phase 6: Documentation (1 week)

User guide creation
Code documentation
Architecture diagrams

Technical Implementation



Initial Pulumi Setup/

```
|-- Pulumi.yaml          # Pulumi project and stack metadata
|-- Pulumi.dev.yaml      # Dev stack settings (config...)
|-- configs/             # Stack-specific configuration files
|   |-- dev.json          # Development stack resources
|   |-- prod.json          # Production stack resources
|   |-- test.json          # Test stack resources
|-- requirements.txt      # Python dependencies
|-- venv/                 # Python virtual environment
|   |-- .env.example        # Template for environment variables
|   |-- .env.local           # Your personal environment variables (create this)
|-- core/                 # Base_resource, factory, resolver
|-- loaders/              # JSON loader and .env reader
|-- resources/             # Python classes for each OCI resource type
`-- main.py                # Entrypoint: one-loop provisioning logic
```

JSON Configuration Example

infrastructure.json

```
{  
  "name": "production-vcn",  
  "cidrBlock": "10.0.0.0/16",  
  "internetGateways": [...]  
  "securityLists": [...]  
  "routeTables": [...]  
  "subnets": [...]  
}
```

Name-Based References: Resources reference each other by human-readable names

Automatic Resolution: Framework resolves names to OCIDs automatically

Core Framework Components

BaseResource (Abstract)

File: core/base_resource.py (21 lines)

Purpose: Abstract base class defining resource interface

Methods:

- create() - Factory method
- build() - Abstract (must implement)

ResourceFactory

File: core/resource_factory.py (38 lines)

Purpose: Factory pattern for dynamic instantiation

Registry: Maps type strings → Python classes

7 resource types registered

DependencyResolver

File: core/dependency_resolver.py (14 lines)

Purpose: Resource reference manager

Methods:

- register_resource(name, resource)
- get_resource_id(name)

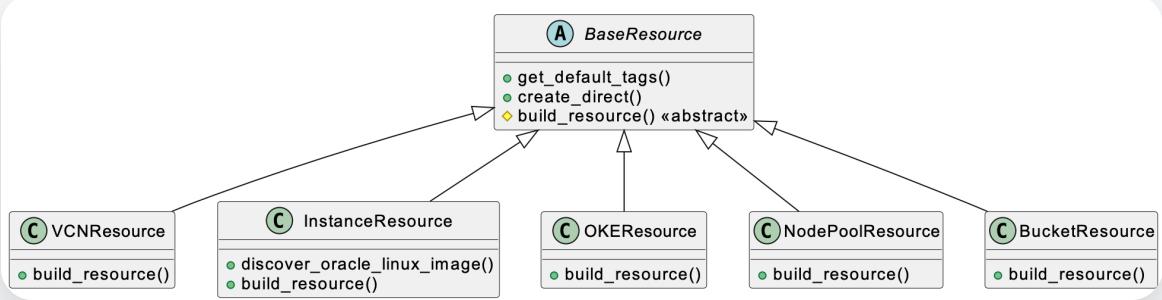
ConfigLoader

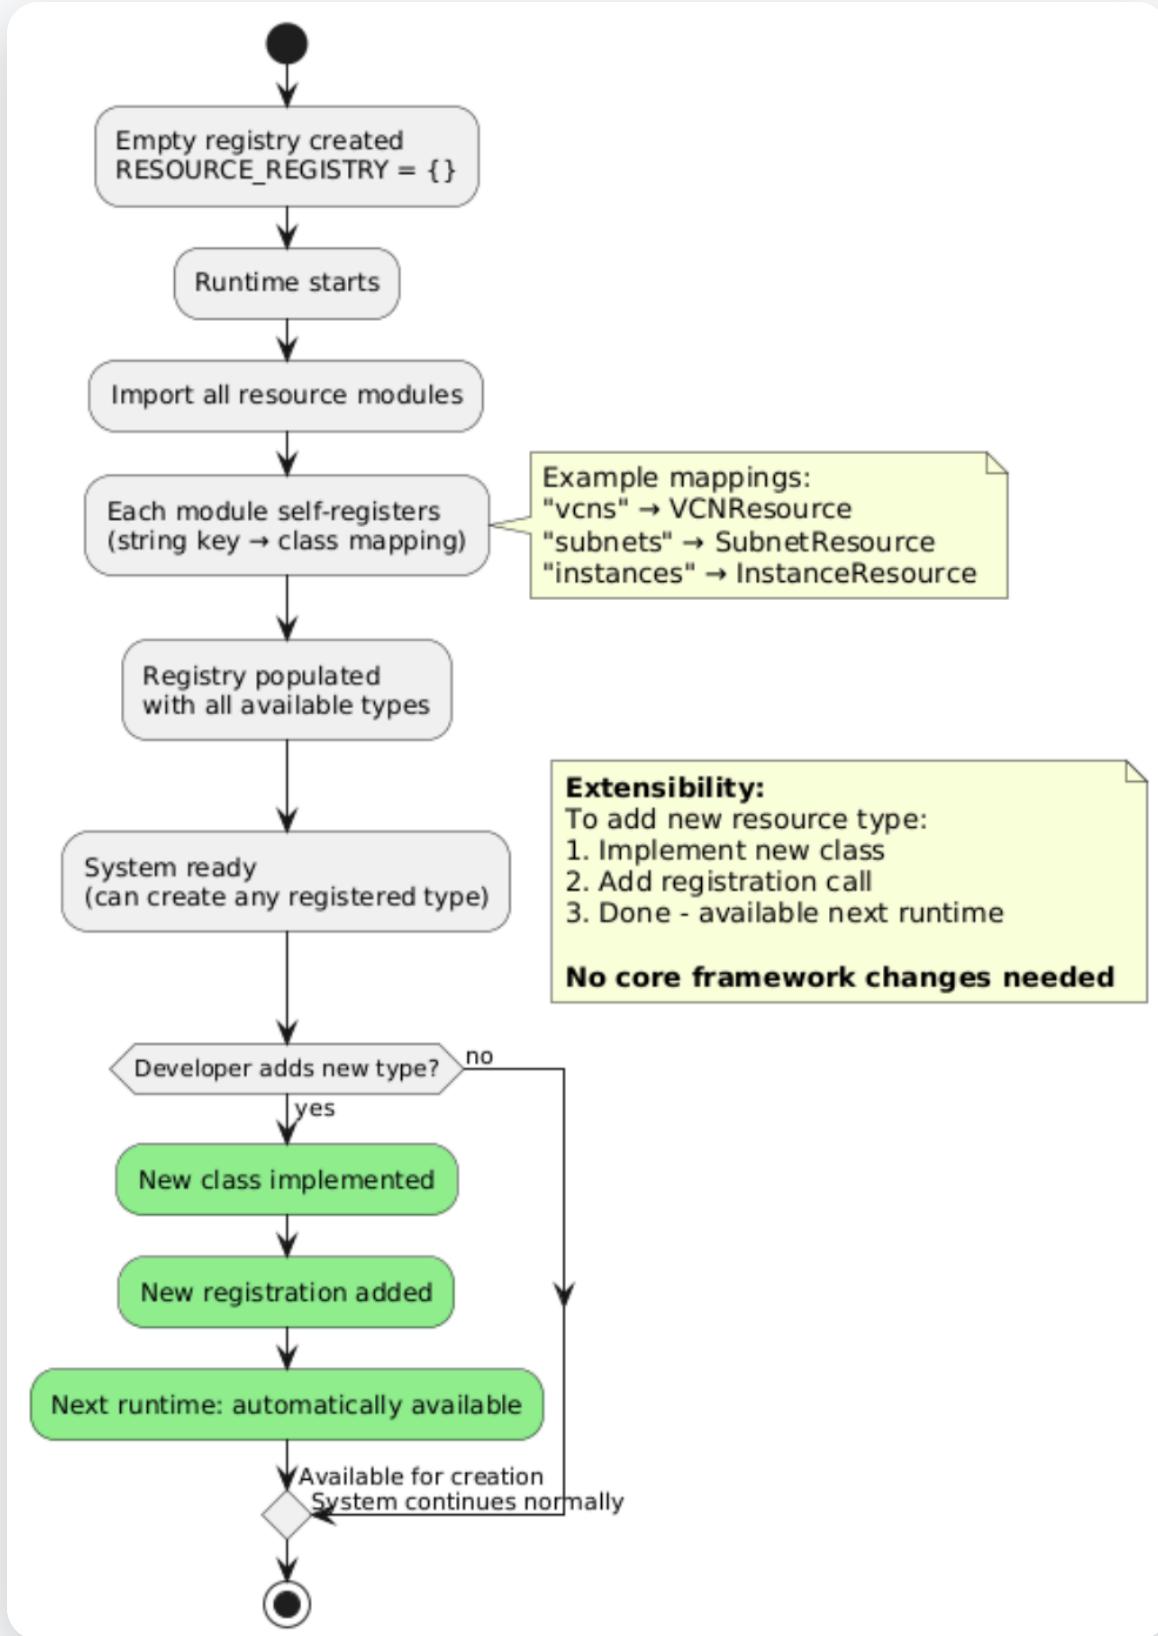
File: config/config_loader.py (25 lines)

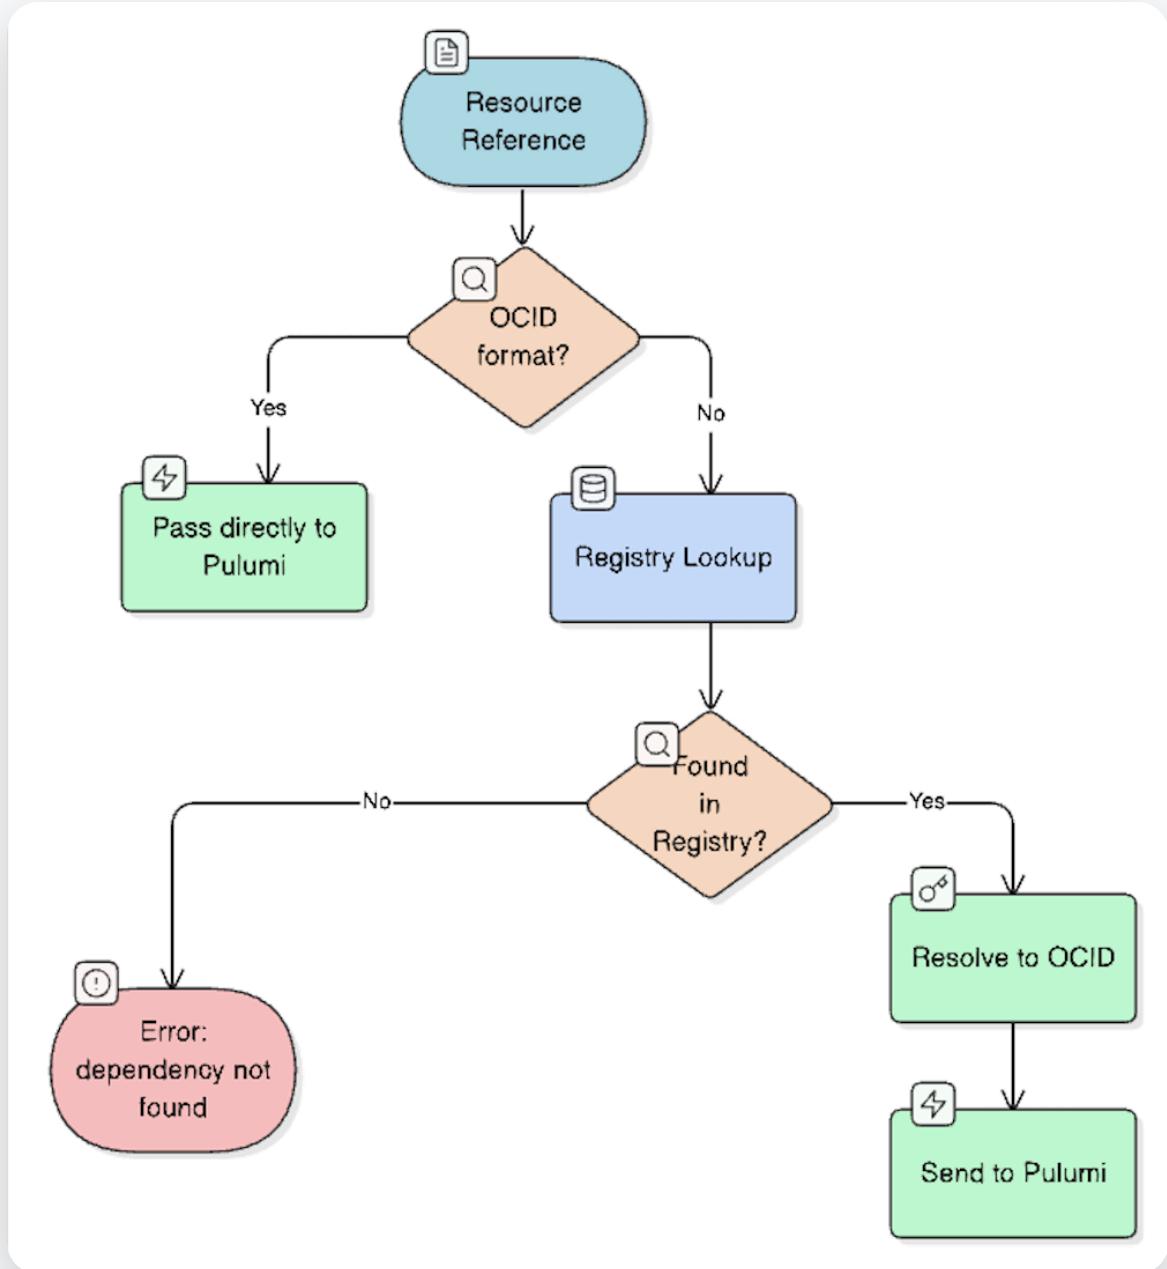
Purpose: Parse infrastructure.json

Methods:

- get_vcns_config()
- get_instances_config()
- get_image_config()







Implemented Resource Types

1. VCN (Virtual Cloud Network)

Purpose: Foundation network container

2. Internet Gateway

Purpose: Enable internet connectivity

3. Security List

Purpose: Firewall rules (ingress/egress)

4. Route Table

Purpose: Network routing configuration

5. Subnet

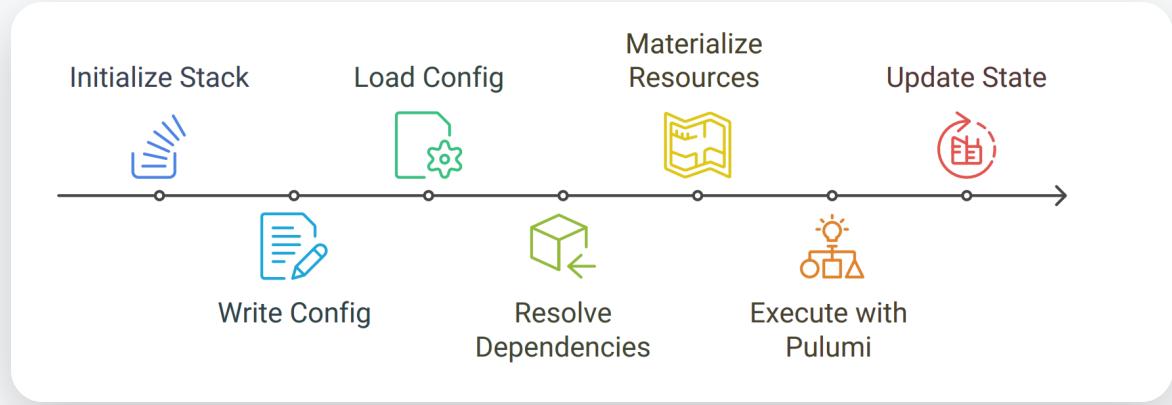
Purpose: Network segment within VCN

6. Compute Instance

Purpose: Virtual machine

7. Object Storage Bucket

Purpose: S3-compatible object storage



use firendly names for Jobs

Passed [REDACTED] created pipeline for commit [REDACTED] 21 hours ago, finished 21 hours ago

1 related merge request: [!2 kept the configs folder](#)

latest ⏪ 2 jobs ⏳ 50 seconds, queued for 15 seconds

Pipeline Jobs 2 Tests 0

Group jobs by Stage Job dependencies

plan

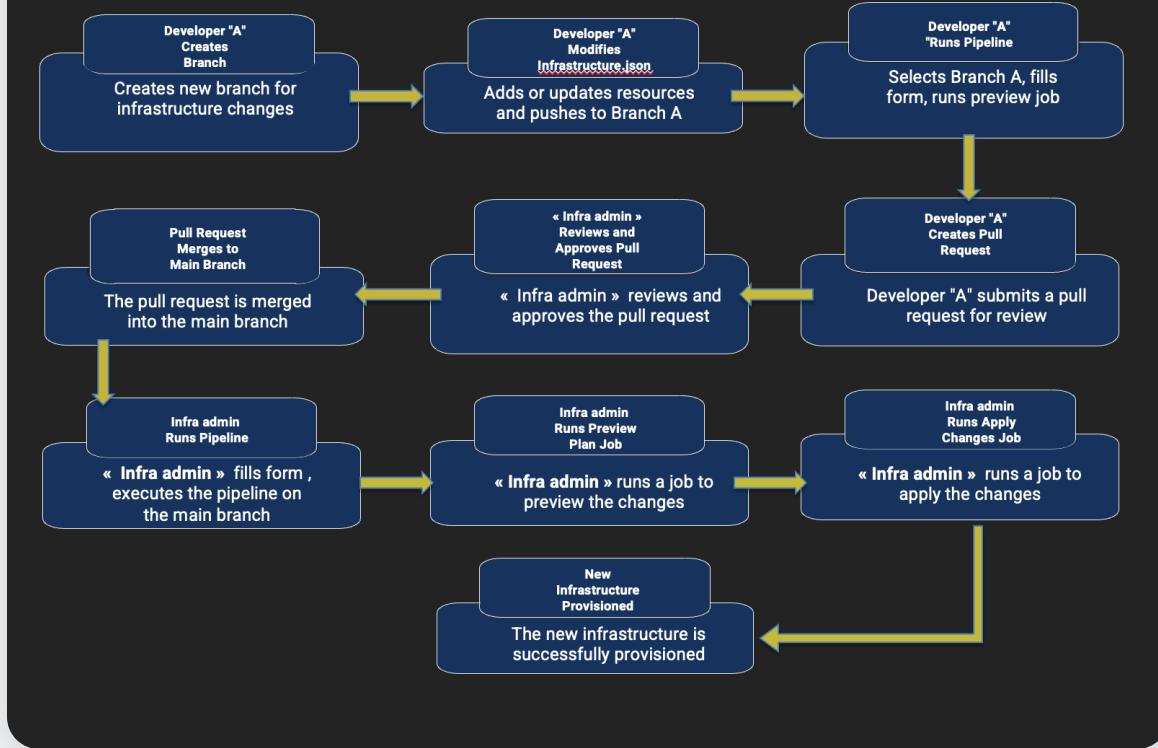
Passed Graal CI Staging Plan

deploy

Passed Graal CI Staging Apply



Infrastructure Change Management Process



Results & Validations

Implementation Results



Deliverables Completed

- Python framework (800 LOC)
- 7 OCI resource types
- JSON configuration system
- GitLab CI/CD pipeline
- Documentation & diagrams



Requirements Met

- ✓ Declarative JSON config
- ✓ Dependency resolution
- ✓ CI/CD integration
- ✓ Multi-environment support
- ✓ State management



Production Ready

- Error handling implemented
- Tested on dev & staging
- Documentation complete
- Team training delivered
- Now in active use by RISQ team

```

└─ pulumi up
  Previewing update (graalci_prod):
    Type          Name      Plan
  + pulumi:pulumi:Stack  graal-iac-graalci_prod  create
  +   OCI::Core:Vcn     production-vcn        create
  +     OCI::Core:RouteTable prod-rt            create
  +     OCI::Core:Subnet   prod-subnet         create
  +     OCI::Core:InternetGateway prod-igw           create
  +     OCI::Core:SecurityList prod-security       create
  +     OCI::ObjectStorage:Bucket prod-data-bucket  create
  +     OCI::Core:Instance    prod-app           create
  Outputs:
    prod-app_public_ip: [unknown]
  Resources:
    + 8 to create

Do you want to perform this update? yes
Updating (graalci_prod):
  Type          Name      Status
  + pulumi:pulumi:Stack  graal-iac-graalci_prod  created (46s)
  +   OCI::ObjectStorage:Bucket prod-data-bucket  created (0.60s)
  +     OCI::Core:Vcn     production-vcn        created (1s)
  +       OCI::Core:RouteTable prod-rt            created (1s)
  +       OCI::Core:InternetGateway prod-igw           created (1s)
  +       OCI::Core:SecurityList prod-security       created (1s)
  +       OCI::Core:Subnet   prod-subnet         created (2s)
  +       OCI::Core:Instance  prod-app           created (36s)
  Outputs:
    prod-app_public_ip: "152.70.52.111"
  Resources:
    + 8 created
  Duration: 48s

└─ ssh -i ~/.ssh/oci_key opc@152.70.52.111
The authenticity of host '152.70.52.111 (152.70.52.111)' can't be established.
ED25519 key fingerprint is SHA256:Yz4gYLSbkaL5AKMLxz0T6PW2yCer0vo4BXbmcHdE6zE.
This key is not known by any other names.
Are you sure you want to continue connecting (yes/no/[fingerprint])? yes
Warning: Permanently added '152.70.52.111' (ED25519) to the list of known hosts.
Activate the web console with: systemctl enable --now cockpit.socket

[opc@ProductionAppServer ~]$ uname -a
Linux ProductionAppServer 5.15.0-312.187.5.1.el8uek.x86_64 #2 SMP Mon Sep 8 12:49:36 PDT 2025 x86_64 x86_64 x86_64 GNU/Linux
[opc@ProductionAppServer ~]$ exit
logout
Connection to 152.70.52.111 closed.

```

Testing & Validation Results

Unit Testing

Scope: Individual resource classes

- Factory pattern registration
- Dependency resolver lookup
- Config loader parsing

✓ All components working correctly

CI/CD Pipeline Testing

Validated:

- Plan stage (preview changes)
- Deploy stage (apply changes)
- Manual approval gate
- State persistence

✓ Pipeline stable over 20+ runs

Integration Testing

Test Cases:

- Single VCN + 1 instance
- 2 VCNs + 2 instances
- Complex network (IGW + RT + SL)
- Bucket creation

✓ End-to-end provisioning successful

Edge Cases Tested

- ✓ Missing dependency (error caught)
- ✓ Invalid CIDR block (validation)
- ✓ Duplicate resource names (detected)
- ✓ Image not compatible with shape

Performance Metrics & Improvements

95%

Time Reduction

From 2-3 hours manual setup
To <5 minutes automated

100%

Consistency

Zero configuration drift
All envs from same templates

85%

Error Reduction

Automated validation
Dependency checking

Before (Manual)

Average setup: 2.5 hours
Error rate: ~15%
Documentation: Often outdated
Reproducibility: Low

After (Automated)

Average setup: 4 minutes
Error rate: <2%
Documentation: Self-documenting JSON
Reproducibility: 100%

Business Impact & Team Benefits

Developer Productivity: Team members can now focus on GraalVM development instead of infrastructure management - estimated 10+ hours saved per week across the team

Faster Testing Cycles: Spin up test environments in minutes instead of hours - enables rapid experimentation and bug fixes

Onboarding Acceleration: New team members can provision infrastructure on day one without specialized training - reduces onboarding time from days to hours

Cost Optimization: All infrastructure tracked in Git - easier to identify and delete unused resources, preventing budget overruns

Future Enhancements

Future Improvements

More Resource Coverage

Expand framework to support additional OCI resources like Load Balancers, Databases, and File Storage

Schema Validation

Implement JSON schema validation to catch configuration errors before deployment

Role Identification

Define clear roles within the team on how to use the framework and establish best practices

Future Work & Enhancement Roadmap

Q1 2026: Extended Resources

Priority: High

- Add Load Balancer support
- Autonomous Database provisioning
- File Storage Service integration
- Network Load Balancer

Effort: 4 weeks

Q2 2026: Validation & Safety

Priority: High

- JSON schema validation
- Pre-deployment cost estimation
- Security policy enforcement
- CIDR conflict detection

Effort: 3 weeks

Q3 2026: Operational Features

Priority: Medium

- Automated drift detection
- Scheduled state refresh jobs
- Notification integration (Slack)
- Resource tagging automation

Effort: 5 weeks

Q4 2026: Multi-Cloud Expansion

Priority: Low

- AWS provider support
- Azure provider support
- Provider abstraction layer
- Cross-cloud networking

Effort: 8 weeks

Lessons Learned & Best Practices

Technical Insights

Start Simple: MVP with 2 resources, then expand - avoid over-engineering upfront

Embrace Patterns: Factory + Resolver patterns made code extensible and maintainable

Test Early: Integration testing caught dependency issues before production

Document as You Go: Inline comments and diagrams saved time later

Challenges Overcome

Dependency Hell: Solved with centralized resolver pattern

State Management: Pulumi's built-in state backend simplified this

Error Handling: Graceful failures with clear messages improved UX

CI/CD Integration: Manual approval gate provided safety net

Process Insights

Agile Works: 2-week sprints with demos kept momentum high

Feedback Loop: Weekly mentor sessions caught issues early

User-Centric: Involving team in design ensured adoption

Advice for Similar Projects

- ✓ Prototype quickly, iterate based on feedback
- ✓ Invest in good abstractions early
- ✓ Automate testing from day one
- ✓ Keep code modular (<100 LOC per file)

Conclusion

Successfully delivered a production-ready IaC framework that transforms infrastructure management for the GraalVM RISQ team

Objectives Achieved

- ✓ JSON-driven configuration
- ✓ 7 OCI resource types
- ✓ GitLab CI/CD integration
- ✓ 95% time reduction
- ✓ Production deployment

Innovation Highlights

- Name-based dependency resolution
- Factory pattern for extensibility
- Declarative JSON interface
- Two-stage approval workflow
- Self-documenting architecture

Measurable Impact

- 10+ hours saved per week
- Zero configuration drift
- 85% error reduction
- 100% reproducibility
- Active team adoption

Future Potential

- Expand to 15+ resource types
- Multi-cloud support
- Advanced validation
- Cost optimization features
- Enterprise-wide adoption

Acknowledgments

Special Thanks

Hamza GHAISI - Oracle Labs Mentor & Technical Supervisor

For invaluable guidance, code reviews, and architectural insights throughout the project

Academic Supervisor

Pr. Hatim LECHGAR

École Hassania des Travaux Publics

For academic guidance and project oversight

GraalVM RISQ Team

For welcoming me to the team, providing feedback, and actively using the framework

Oracle Labs & EHTP

For providing this incredible learning opportunity and fostering innovation in infrastructure automation

References & Resources

Technical Documentation

Pulumi Documentation
pulumi.com/docs

Oracle Cloud Infrastructure API Reference
docs.oracle.com/iaas/api

GitLab CI/CD Documentation
docs.gitlab.com/ee/ci

Python Design Patterns
refactoring.guru/design-patterns

Key Technologies

Pulumi: v3.0.0+ (IaC framework)

Python: 3.13 (programming language)

OCI Provider: v2.0.0+ (cloud provider)

GitLab CI/CD: (automation platform)

JSON: (configuration format)

Project Resources

GitHub Repository:
github.com/y-dbaichi/pfe-pulumi
Internal Confluence Documentation
Architecture Diagrams (25 PNG files)

Contact Information

Yassine DBAICHI

Infrastructure Automation Engineer
Intern

Oracle Labs - GraalVM RISQ Team
École Hassania des Travaux Publics

Thank You

Questions & Discussion

Yassine DBAICHI

Infrastructure as Code Support in Graal CI

Oracle Labs - GraalVM RISQ Team

Academic Year 2024-2025

